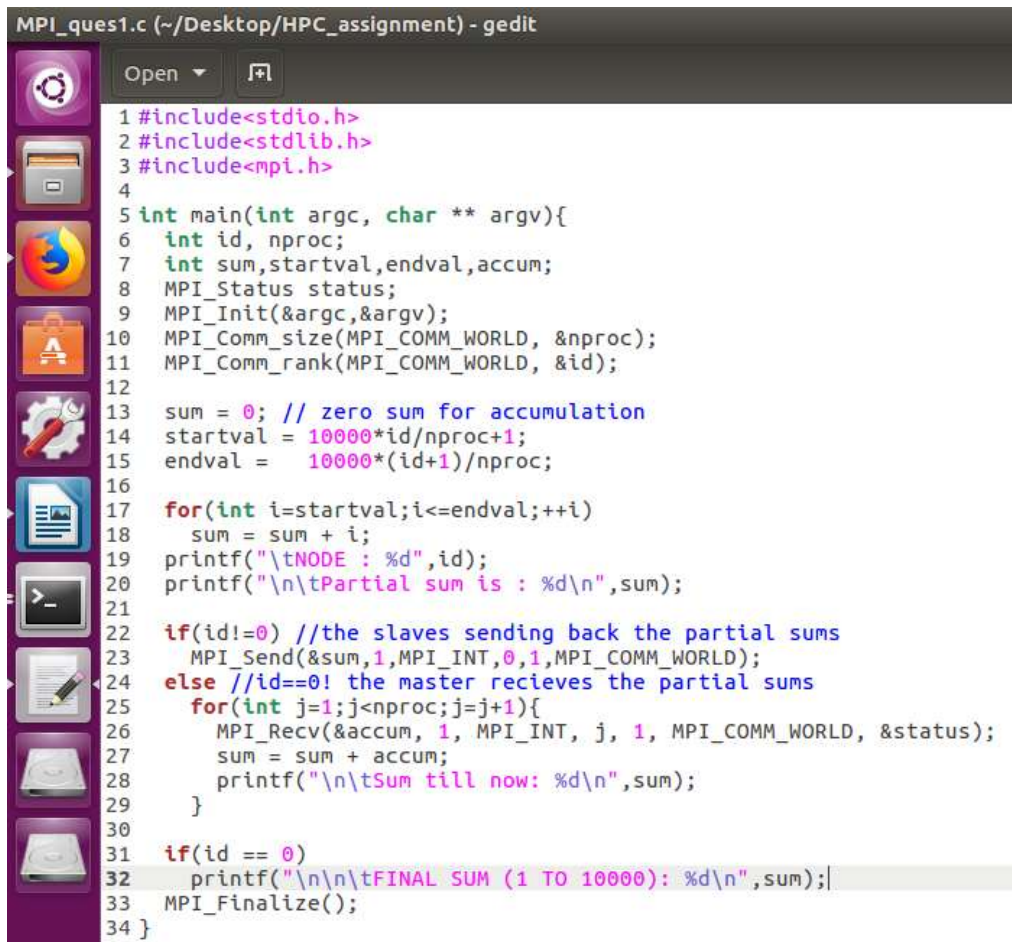


HIGH PERFORMANCE COMPUTING

ASSIGNMENT

1. Write a program to sum the numbers from 1 to 10000 in sequential program. To parallelize this procedure using MPI, divide the process into sub-processes and write the code. The code should assign different numbers to be summed to different processes, and outputs the sum and sub-sums in different processes.

Code:



```

MPI_ques1.c (~/Desktop/HPC_assignment) - gedit
1 #include<stdio.h>
2 #include<stdlib.h>
3 #include<mpi.h>
4
5 int main(int argc, char ** argv){
6     int id, nproc;
7     int sum, startval, endval, accum;
8     MPI_Status status;
9     MPI_Init(&argc, &argv);
10    MPI_Comm_size(MPI_COMM_WORLD, &nproc);
11    MPI_Comm_rank(MPI_COMM_WORLD, &id);
12
13    sum = 0; // zero sum for accumulation
14    startval = 10000*id/nproc+1;
15    endval = 10000*(id+1)/nproc;
16
17    for(int i=startval; i<=endval; ++i)
18        sum = sum + i;
19    printf("\tNODE : %d", id);
20    printf("\n\tPartial sum is : %d\n", sum);
21
22    if(id!=0) //the slaves sending back the partial sums
23        MPI_Send(&sum, 1, MPI_INT, 0, 1, MPI_COMM_WORLD);
24    else //id==0! the master receives the partial sums
25        for(int j=1; j<nproc; j=j+1){
26            MPI_Recv(&accum, 1, MPI_INT, j, 1, MPI_COMM_WORLD, &status);
27            sum = sum + accum;
28            printf("\n\tSum till now: %d\n", sum);
29        }
30
31    if(id == 0)
32        printf("\n\n\tFINAL SUM (1 TO 10000): %d\n", sum);
33    MPI_Finalize();
34 }
  
```

Output:

```

dhruv@dhruv-Inspiron-5559:~/Desktop/HPC_assignment$ mpicc MPI_ques1.c
dhruv@dhruv-Inspiron-5559:~/Desktop/HPC_assignment$ mpirun -np 4 ./a.out
NODE : 0
Partial sum is : 3126250

Sum till now: 12502500

Sum till now: 28128750

Sum till now: 50005000

FINAL SUM (1 TO 10000): 50005000
NODE : 1
Partial sum is : 9376250
NODE : 2
Partial sum is : 15626250
NODE : 3
Partial sum is : 21876250
  
```



```

MPI_ques2.c (~/Desktop/HPC_assignment) - gedit
73 }
74 }
75 }
76 /* Convert to half the round-trip time */
77 tmin = tmin / 2.0;
78 if (rank == 0) {
79     printf( "Vector\t%d\t%d\t%f\t%f\n", n, stride, tmin, n * sizeof(double) * 1.0e-6 / tmin );
80 }
81 MPI_Type_free( &vec1 );
82
83 /* Use a variable vector type */
84 blocklens[0] = 1;
85 blocklens[1] = 1;
86 indices[0] = 0;
87 indices[1] = stride * sizeof(double);
88 old_types[0] = MPI_DOUBLE;
89 old_types[1] = MPI_UB;
90 MPI_Type_struct( 2, blocklens, indices, old_types, &vec_n );
91 MPI_Type_commit( &vec_n );
92
93 tmin = 1000;
94 for (k=0; k<NUMBER_OF_TESTS; k++) {
95     if (rank == 0) {
96         /* Make sure both processes are ready */
97         MPI_Sendrecv( MPI_BOTTOM, 0, MPI_INT, 1, 14,
98                     MPI_BOTTOM, 0, MPI_INT, 1, 14, MPI_COMM_WORLD,
99                     &status );
100         t1 = MPI_Wtime();
101         for (j=0; j<nloop; j++) {
102             MPI_Send( buf, n, vec_n, 1, k, MPI_COMM_WORLD );
103             MPI_Recv( buf, n, vec_n, 1, k, MPI_COMM_WORLD, &status );
104         }
105         t2 = (MPI_Wtime() - t1) / nloop;
106         if (t2 < tmin) tmin = t2;
107     }
108     else if (rank == 1) {
109         /* Make sure both processes are ready */
110         MPI_Sendrecv( MPI_BOTTOM, 0, MPI_INT, 0, 14,
111                     MPI_BOTTOM, 0, MPI_INT, 0, 14, MPI_COMM_WORLD,
112                     &status );

```

```

File Edit View Search Tools Documents Help
113     for (j=0; j<nloop; j++) {
114         MPI_Recv( buf, n, vec_n, 0, k, MPI_COMM_WORLD, &status );
115         MPI_Send( buf, n, vec_n, 0, k, MPI_COMM_WORLD );
116     }
117 }
118 }
119 /* Convert to half the round-trip time */
120 tmin = tmin / 2.0;
121 if (rank == 0) {
122     printf( "Struct\t%d\t%d\t%f\t%f\n", n, stride, tmin, n * sizeof(double) * 1.0e-6 / tmin );
123 }
124
125 MPI_Type_free( &vec_n );
126
127 /* Use user-packing with known stride */
128 tmin = 1000;
129 for (k=0; k<NUMBER_OF_TESTS; k++) {
130     if (rank == 0) {
131         /* Make sure both processes are ready */
132         MPI_Sendrecv( MPI_BOTTOM, 0, MPI_INT, 1, 14,
133                     MPI_BOTTOM, 0, MPI_INT, 1, 14, MPI_COMM_WORLD,
134                     &status );
135         t1 = MPI_Wtime();
136         for (j=0; j<nloop; j++) {
137             /* If the compiler isn't good at unrolling and changing
138             multiplication to indexing, this won't be as good as
139             it could be */
140             for (i=0; i<n; i++)
141                 lbuf[i] = buf[i*stride];
142             MPI_Send( lbuf, n, MPI_DOUBLE, 1, k, MPI_COMM_WORLD );
143             MPI_Recv( lbuf, n, MPI_DOUBLE, 1, k, MPI_COMM_WORLD, &status );
144             for (i=0; i<n; i++)
145                 buf[i*stride] = lbuf[i];
146         }
147         t2 = (MPI_Wtime() - t1) / nloop;
148         if (t2 < tmin) tmin = t2;
149     }
150     else if (rank == 1) {
151         /* Make sure both processes are ready */
152         MPI_Sendrecv( MPI_BOTTOM, 0, MPI_INT, 0, 14,

```



```

MPI_ques2.c (~/Desktop/HPC_assignment) - gedit
Open  [icon]

150 else if (rank == 1) {
151     /* Make sure both processes are ready */
152     MPI_Sendrecv( MPI_BOTTOM, 0, MPI_INT, 0, 14,
153                 MPI_BOTTOM, 0, MPI_INT, 0, 14, MPI_COMM_WORLD,
154                 &status );
155     for (j=0; j<nloop; j++) {
156         MPI_Recv( lbuf, n, MPI_DOUBLE, 0, k, MPI_COMM_WORLD, &status );
157         for (i=0; i<n; i++)
158             buf[i*stride] = lbuf[i];
159         for (i=0; i<n; i++)
160             lbuf[i] = buf[i*stride];
161         MPI_Send( lbuf, n, MPI_DOUBLE, 0, k, MPI_COMM_WORLD );
162     }
163 }
164 }
165 /* Convert to half the round-trip time */
166 tmin = tmin / 2.0;
167 if (rank == 0) {
168     printf( "User\t%d\t%d\t%f\t%f\n", n, stride, tmin, n * sizeof(double) * 1.0e-6 / tmin );
169 }
170 }
171 /* Use user-packing with known stride, using addition in the user
172    copy code */
173 tmin = 1000;
174 for (k=0; k<NUMBER_OF_TESTS; k++) {
175     if (rank == 0) {
176         /* Make sure both processes are ready */
177         MPI_Sendrecv( MPI_BOTTOM, 0, MPI_INT, 1, 14,
178                     MPI_BOTTOM, 0, MPI_INT, 1, 14, MPI_COMM_WORLD,
179                     &status );
180         t1 = MPI_Wtime();
181         for (j=0; j<nloop; j++) {
182             /* If the compiler isn't good at unrolling and changing
183                multiplication to indexing, this won't be as good as
184                it could be */
185             in_p = buf; out_p = lbuf;
186             for (i=0; i<n; i++) {
187                 out_p[i] = *in_p; in_p += stride;
188             }
189         }
190     }
191     MPI_Send( lbuf, n, MPI_DOUBLE, 1, k, MPI_COMM_WORLD );
192     MPI_Recv( lbuf, n, MPI_DOUBLE, 1, k, MPI_COMM_WORLD, &status );
193     out_p = buf; in_p = lbuf;
194     for (i=0; i<n; i++) {
195         *out_p = in_p[i]; out_p += stride;
196     }
197     t2 = (MPI_Wtime() - t1) / nloop;
198     if (t2 < tmin) tmin = t2;
199 }
200 else if (rank == 1) {
201     /* Make sure both processes are ready */
202     MPI_Sendrecv( MPI_BOTTOM, 0, MPI_INT, 0, 14,
203                 MPI_BOTTOM, 0, MPI_INT, 0, 14, MPI_COMM_WORLD,
204                 &status );
205     for (j=0; j<nloop; j++) {
206         MPI_Recv( lbuf, n, MPI_DOUBLE, 0, k, MPI_COMM_WORLD, &status );
207         in_p = lbuf; out_p = buf;
208         for (i=0; i<n; i++) {
209             *out_p = in_p[i]; out_p += stride;
210         }
211         out_p = lbuf; in_p = buf;
212         for (i=0; i<n; i++) {
213             out_p[i] = *in_p; in_p += stride;
214         }
215     }
216     MPI_Send( lbuf, n, MPI_DOUBLE, 0, k, MPI_COMM_WORLD );
217 }
218 }
219 /* Convert to half the round-trip time */
220 tmin = tmin / 2.0;
221 if (rank == 0) {
222     printf( "User(odd)\t%d\t%d\t%f\t%f\n",
223           n, stride, tmin, n * sizeof(double) * 1.0e-6 / tmin );
224 }
225 MPI_Finalize( );
226 return 0;
227 }
228 }

```

```

MPI_ques2.c (~/Desktop/HPC_assignment) - gedit
Open  [icon]

189
190 MPI_Send( lbuf, n, MPI_DOUBLE, 1, k, MPI_COMM_WORLD );
191 MPI_Recv( lbuf, n, MPI_DOUBLE, 1, k, MPI_COMM_WORLD, &status );
192 out_p = buf; in_p = lbuf;
193 for (i=0; i<n; i++) {
194     *out_p = in_p[i]; out_p += stride;
195 }
196 }
197 t2 = (MPI_Wtime() - t1) / nloop;
198 if (t2 < tmin) tmin = t2;
199 }
200 else if (rank == 1) {
201     /* Make sure both processes are ready */
202     MPI_Sendrecv( MPI_BOTTOM, 0, MPI_INT, 0, 14,
203                 MPI_BOTTOM, 0, MPI_INT, 0, 14, MPI_COMM_WORLD,
204                 &status );
205     for (j=0; j<nloop; j++) {
206         MPI_Recv( lbuf, n, MPI_DOUBLE, 0, k, MPI_COMM_WORLD, &status );
207         in_p = lbuf; out_p = buf;
208         for (i=0; i<n; i++) {
209             *out_p = in_p[i]; out_p += stride;
210         }
211         out_p = lbuf; in_p = buf;
212         for (i=0; i<n; i++) {
213             out_p[i] = *in_p; in_p += stride;
214         }
215     }
216     MPI_Send( lbuf, n, MPI_DOUBLE, 0, k, MPI_COMM_WORLD );
217 }
218 }
219 /* Convert to half the round-trip time */
220 tmin = tmin / 2.0;
221 if (rank == 0) {
222     printf( "User(odd)\t%d\t%d\t%f\t%f\n",
223           n, stride, tmin, n * sizeof(double) * 1.0e-6 / tmin );
224 }
225 MPI_Finalize( );
226 return 0;
227 }
228 }

```

Output:

```
dhruv@dhruv-Inspiron-5559: ~/Desktop/HPC_assignment
dhruv@dhruv-Inspiron-5559:~/Desktop/HPC_assignment$ mpicc MPI_ques2.c
dhruv@dhruv-Inspiron-5559:~/Desktop/HPC_assignment$ mpirun -np 4 ./a.out
Kind      n      stride  time (sec)    Rate (MB/sec)
Vector    1000    24      0.000005      1673.955201
Struct    1000    24      0.000005      1751.732289
User      1000    24      0.000012      683.389654
User(add)      1000    24      0.000009      867.151622
dhruv@dhruv-Inspiron-5559:~/Desktop/HPC_assignment$
```